CS 2110 Homework 7
Recursive Assembly

Daniel Becker, Maddie Brickell, Michael Xu, and Cliff Panos

Fall 2018

Contents

1 Overview

In this assignment, you will be implementing several subroutines in assembly, taking advantage of the calling convention that has been outlined in class/lecitation. **For each subroutine, you will be required to implement the entire calling convention. This includes the buildup and teardown on the stack**.

2 Instructions

You have been provided three assembly files, **string_ops.asm, handshake.asm, and buildheap.asm**. These files contain a label for each of the subroutines that you will be asked to implement. **DO NOT RENAME THE LABELS!** You can add more to suit your needs, but don't change the existing ones. Implement the following subroutines, use complx to debug them, and submit your files on Gradescope.

It is in your best interest to do the problems in order. The last one builds on the first two.

2.1 Part 1: String Operations (20 points)

The first part of this assignment is to implement a method to count the occurrences of a character in a string. This will utilize the `strlen` subroutine, which has been implemented for you.

2.1.1 String Length

To help you get started on the homework, we have provided you with the implementation of `strlen`. It includes the buildup of the stack, body of the function, and teardown of the stack. The pseudocode is provided here for your reference, but **you do not have to implement this yourself**. You will, however, have to call it from the next subroutine, so make sure you are familiar with how it works.

Pseudocode

```
int strlen(String s) {
    int count = 0;
    while (s.charAt(count) != "\0") {
        count++;
    }
}
```

2.1.2 Count Character

Complete the `count_occurrence` function in the `string_ops.asm` file. This function should count how many times a given character appears in a string. This method should call the `strlen` method that you implemented in the previous part.

Pseudocode

```
int count_occurrence(String s, char c) {
    int count = 0;
    for (int i = 0; i < strlen(s); i++) {
        if (s.charAt(i) == c) {
            count++;
        }
```

```
        }
    return count;
}
```

2.2 Part 2: Handshaking (30 points)

The second part of this assignment is solving the handshake problem.

Background

With n people at a party, find the maximum number of handshakes possible given that any two people can only shake each other's hand once. This problem can be easily solved using recursion and the pseudocode for this problem is in the next section!

Pseudocode

```
int handshake(int n) {
    if (n == 0)
        return 0;
    else
        return (n - 1) + handshake(n - 1);
}
```

2.3 Part 3: Build Heap (50 points)

Background

If you have taken CS 1332 (don't worry if you haven't), then you might be familiar with the data structure known as the binary heap. A binary heap (**specifically a max heap**) is essentially a binary tree with the property that for every node, its children are less than or equal to it. This means that the root node of the heap is the largest element. The cool thing about binary heaps is that they can be represented as arrays, where a node at index **n** has its left child at index **2n + 1** and its right child at index **2n + 2**. The goal of the next two problems will be to implement a function that will order an array of elements so that they form a valid heap. While it may help you to know more about binary heaps in order to understand what is going on (you can ready about them at `https://www.geeksforgeeks.org/binary-heap/`), you don't actually have to know anything about them since we give you the pseudocode.

2.3.1 Heapify (35 points)

The first subroutine that we will have you implement is heapify. Heapify takes in an array, its length, and a given node i and then ensures that the heap property of a node being larger than its children is maintained for that node. If you have to swap two elements, then heapify recursively re-checks the heap property for child nodes. Once we have implemented, we will be able to call it for every node to actually convert a given array into a heap.

Your job is to implement heapify in the `buildheap.asm` file, starting at the `heapify` label.

Heapify Pseudocode

```
void heapify(int[] arr, int n, int i) {
    // The following are all indices, not values
```

```
    int largest = i; // Initialize largest as root
    int leftChild = 2*i + 1; // left = 2*i + 1
    int rightChild = 2*i + 2; // right = 2*i + 2

    // If left child is larger than root
    if (leftChild < n && arr[leftChild] > arr[largest])
        largest = leftChild;

    // If right child is larger than largest so far
    if (rightChild < n && arr[rightChild] > arr[largest])
        largest = rightChild;

    // If largest is not root
    if (largest != i)
    {
        swap(arr[i], arr[largest]);

        // Recursively heapify the affected sub-tree
        heapify(arr, n, largest);
    }
}
```

2.3.2 Build Heap (15 points)

Once you've implemented heapify, buildheap is easy. All you have to do is start at the end of the array and iterate backwards over each element, calling heapify on each one.

In `buildheap.asm`, implement the following pseudocode, starting at the `buildheap` label.

Build Heap Pseudocode

```
void buildheap(int arr[], int n) {
    for (int i = n; i >= 0; i--)
        heapify(arr, n, i);
}
```

3   Deliverables

Please upload the following files onto the assignment on Gradescope:

1. `string_ops.asm`

2. `handshake.asm`

3. `buildheap.asm`

Be sure to check your score to see if you submitted the right files, as well as your email frequently until the due date of the assignment for any potential updates.

4 LC-3 Assembly Programming Requirements

4.1 Overview

1. Your code must assemble with **NO WARNINGS OR ERRORS**. To assemble your program, open the file with Complx. It will complain if there are any issues. **If your code does not assemble you WILL get a zero for that file.**

2. **Comment your code!** This is especially important in assembly, because it's much harder to interpret what is happening later, and you'll be glad you left yourself notes on what certain instructions are contributing to the code. Comment things like what registers are being used for and what less intuitive lines of code are actually doing. To comment code in LC-3 assembly just type a semicolon (;), and the rest of that line will be a comment.

3. Avoid stating the obvious in your comments, it doesn't help in understanding what the code is doing.

   Good Comment

   ```
   ADD R3, R3, -1          ; counter--
   BRp LOOP                ; if counter == 0 don't loop again
   ```

   Bad Comment

   ```
   ADD R3, R3, -1          ; Decrement R3
   BRp LOOP                ; Branch to LOOP if positive
   ```
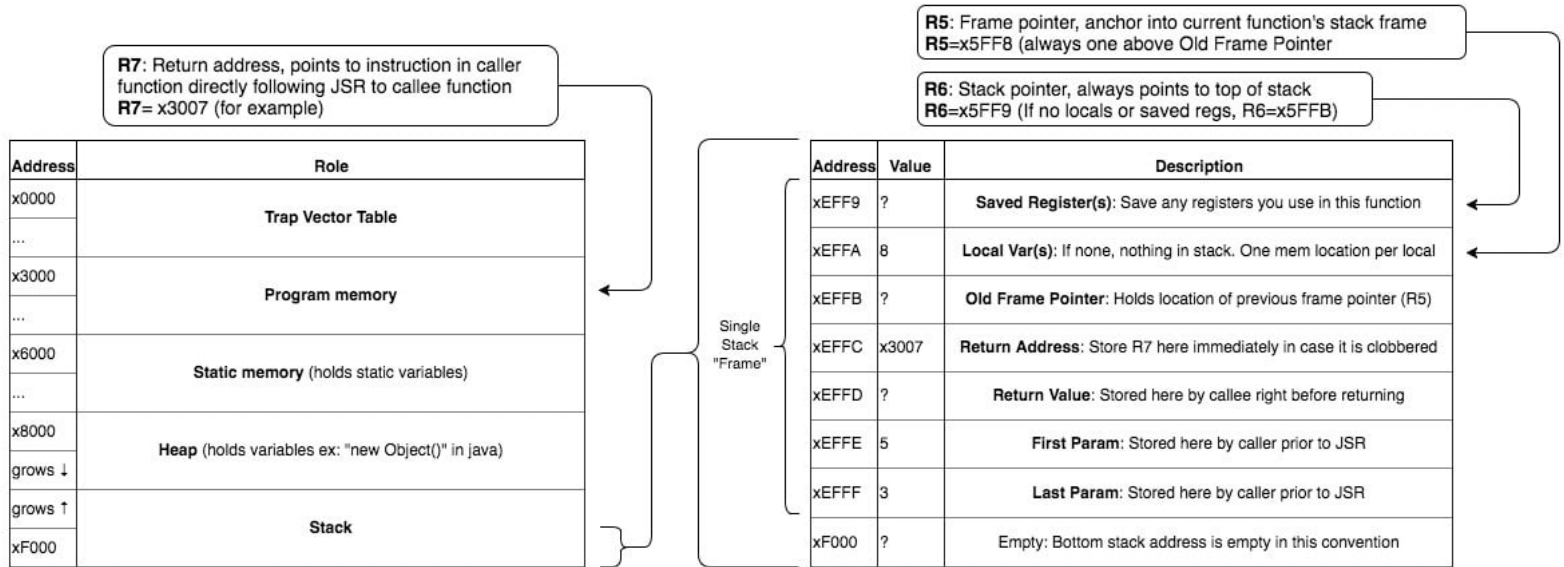
4. **DO NOT assume that ANYTHING in the LC-3 is already zero.** Treat the machine as if your program was loaded into a machine with random values stored in the memory and register file.

5. Following from 3. You can randomize the memory and load your program by doing File - Randomize and Load.

6. Use the LC-3 calling convention. This means that all local variables, frame pointer, etc. . . must be pushed onto the stack. Our autograder will be checking for correct stack setup.

7. Start the stack at xF000. **The stack pointer always points to the last used stack location.** This means you will allocate space **first**, then store onto the stack pointer.

8. Do NOT execute any data as if it were an instruction (meaning you should put .fills after **HALT** or RET).

9. Do not add any comments beginning with @plugin or change any comments of this kind.

10. **Test your assembly.** Don't just assume it works and turn it in.

4.2 LC-3 Calling Convention Guide

A handy assembly guide written by Kyle Murray (RIP) follows on the next page:

Overview

On the left is a grand-scale layout of memory in a computer. On the right is a zoomed in view of the stack. The state of the stack is a snapshot during the execution of some function (we will call it "foo"). Since the stack begins at xF000, we know that we are currently executing the first function called by main (main itself does not have its own stack frame). At this state in time, the main function is the "caller", and the currently executing function, foo, is the "callee". Appropriate example addresses and values are included throughout this guide, but note that these values are not necessarily the standard.

R7: Return address, points to instruction in caller function directly following JSR to callee function
R7= x3007 (for example)

R5: Frame pointer, anchor into current function's stack frame
R5=x5FF8 (always one above Old Frame Pointer

R6: Stack pointer, always points to top of stack
R6=x5FF9 (If no locals or saved regs, R6=x5FFB)

Address	Role
x0000	
...	Trap Vector Table
x3000	
...	Program memory
x6000	
...	Static memory (holds static variables)
x8000	
grows ↓	Heap (holds variables ex: "new Object()" in java)
grows ↑	
xF000	Stack

Single Stack "Frame"

Address	Value	Description
xEFF9	?	**Saved Register(s)**: Save any registers you use in this function
xEFFA	8	**Local Var(s)**: If none, nothing in stack. One mem location per local
xEFFB	?	**Old Frame Pointer**: Holds location of previous frame pointer (R5)
xEFFC	x3007	**Return Address**: Store R7 here immediately in case it is clobbered
xEFFD	?	**Return Value**: Stored here by callee right before returning
xEFFE	5	**First Param**: Stored here by caller prior to JSR
xEFFF	3	**Last Param**: Stored here by caller prior to JSR
xF000	?	**Empty**: Bottom stack address is empty in this convention

Walkthrough

In the following walkthrough, function "foo" will call function "bar", and you will be shown snapshots of the stack at different stages of the calling convention. Note that "caller/foo" and "callee/bar" will be used interchangeably. For transitions 2, 3, and 4 (not 1, 2.1, or 3.0), the instructions can be EXACTLY the between different functions, a "cookie cutter" set of instructions. The following pseudocode describes the two functions (they don't do anything important, they are just examples of functions you might have to recreate in assembly):

```
int foo (int a, int b) {
    int local_var = a + b;
    answer = bar(a, local_var);
    return answer;
}
```

```
int bar (int c, int d) {
    int var1 = c + d;
    int var2 = c - d;
    answer = var1 / var2;
    return answer;
}
```

State 0 - During Foo's Execution

```
int foo (int a, int b) {
    int local_var = a + b;
    answer = bar(a, local_var);    ← PC
    return answer;
}
```

The function has stored local_var but not yet called bar. Foo is about to call bar so we will now refer to foo as the "caller" and bar as the "callee". Remember the PC points to the *planned-to-be-next instruction*.

Address	Value	Description
xEFF9	?	**Foo's Saved Register(s)**: Foo saved all registers it used before use
xEFFA	8	**Foo's Local Var(s)**: stores value of local_var
xEFFB	?	**Foo's Old Frame Pointer**: Foo's caller is main so this is garbage
xEFFC	x3007	**Foo's Return Address**: Points to instruction in main to return to
xEFFD	?	**Foo's Return Value**: currently unknown
xEFFE	5	**Foo's First Param**: a
xEFFF	3	**Foo's Last Param**: b
xF000	?	**Empty** (main's frame)

R6 → xEFF9
R5 → xEFFA

Caller Setup (Transition 1)

Before calling JSR BAR, foo must initiate the creation of bar's stack frame by giving bar its params and moving the stack pointer (R6) accordingly. That's it.

Tasks
Tasks directly translate to assembly instructions

Move stack pointer up enough spaces to make room for params.
Store params on stack.
JSR to BAR.

State 1 - Foo to Bar handoff

```
int foo (int a, int b) {
    int local_var = a + b;
    answer = bar(a, local_var);      [PC before JSR
    return answer;                    R7 after JSR]
}
int bar (int c, int d) {
    int var1 = c + d;                [PC after JSR]
    int var2 = c - d;
    answer = var1 / var2;
    return answer;
}
```

This is the snapshot of the stack right as JSR is called. Notice the callee's (bar's) frame has been started but is incomplete. Think of this stage as the handoff between caller and callee. Immediately after JSR is called, R7 has the value of the return address and our program is now executing in bar.

	Address	Value	Description	
Callee's (bar's) stack frame	xEFF7	5	**Callee's First param: c**	← R6
	xEFF8	8	**Callee's Last param: d**	
Caller's (foo's) stack frame	xEFF9	?	**Caller's Saved Register(s):** Foo saved any registers before use	
	xEFFA	8	**Caller's Local Var(s)**	← R5
	xEFFB	?	**Caller's Old Frame Pointer:** Foo's caller is main so this is garbage	
	xEFFC	x3007	**Caller's Return Address:** Points to main	
	xEFFD	?	**Caller's Return Value:** currently unknown	
	xEFFE	5	**Caller's First Param:** a	
	xEFFF	3	**Caller's Last Param:** b	
	xF000	?	**Empty** (main's frame)	

Callee Setup (Transition 2)

As soon as our new function (bar) is called, we have to finish creating the stack frame and store our important values: return address (R7) and old frame pointer (R5). You should use R6 as your reference point.

Tasks

Move stack pointer up to make room for everything.
Store current value of R5 (caller's frame pointer) onto the stack.
Store the return address (currently stored in R7) onto the stack.
Update R5 to point to one above the position of the old frame pointer in the current frame.

Callee Setup cont (Transition 2.1)

We must now make room for local variables and store the values of any registers you plan on using on the stack. They will be restored later, immediately prior to returning to the caller (foo).

Tasks

Move stack pointer up enough to make room for locals and saved registers.
Store registers you plan to use onto stack.
Store value of local variables whenever they are calculated (not part of this transition)

State 2 - During Bar's Execution

```
int foo (int a, int b) {
    int local_var = a + b;
    answer = bar(a, local_var);
    return answer;                    ←  [ Return address ]
}

int bar (int c, int d) {
    int var1 = c + d;                 ←
    int var2 = c - d;                 ←  [ PC ]
    answer = var1 / var2;
    return answer;
}
```

This is the snapshot of the stack once bar has finished setting up its stack frame. It may now execute all of its logic freely, making sure to only use registers R0-R4 if they have been saved on the stack (in this example just R0 and R1). Notice that this state is similar to State 0, except it is for bar instead of foo. Also note that if bar decided to call a function (eg "baz"), bar would become the caller in that relationship and baz would be the callee.

Address	Value	Description
xEFF0	ex: 6	**Bar's Saved Register: R0**
xEFF1	ex: 7	**Bar's Saved Register: R1**
xEFF2	13	**Bar's Last Local Var**: stores value of var2
xEFF3	-3	**Bar's First Local Var**: stores value of var1
xEFF4	**x5FFA**	**Bar's Old Frame Pointer**: points to Foo's old FP/R5
xEFF5	**See Left**	**Bar's Return Address**: Points to instruction in foo to return to
xEFF6	?	**Bar's Return Value**: currently unknown
xEFF7	5	**Bar's First Param**: c
xEFF8	8	**Bar's Last Param**: d
xEFF9	?	**Caller's Saved Register(s)**: Foo saved any registers before use
xEFFA	8	**Caller's Local Var(s)**: int local_var
xEFFB	?	**Caller's Old Frame Pointer**: Foo's caller is main so this is garbage
xEFFC	x3007	**Caller's Return Address**: Points to main
...	...	...

Callee's (bar's) stack frame — xEFF0 to xEFF8.
Caller's (foo's) stack frame — xEFF9 onward.

R6 → xEFF0
R5 → xEFF3

Callee Teardown (Transition 3.0)

Once we have our return value and are ready to return to the caller, we must first store the return value on the stack, then restore any registers we used. You should use R5 as your reference point

Tasks

Store return value onto stack
Restore any registers (R0-R4) you used from stack. In bar's case we restore R0 and R1

Callee Teardown (Transition 3)

Now for the cookie cutter instructions:
Restore R5 (caller's frame pointer), restore R7 (our return address into the caller), pop down R6 to point to the return value (making it easy for the caller to find it), and return.

Tasks

Restore R5 (frame pointer) from the stack
Restore R7 (return address) from the stack
Pop down R6 (stack pointer) to point to the return value
Return

State 3 - Bar hands back to Foo

```
int foo (int a, int b) {
    int local_var = a + b;
    answer = bar(a, local_var);
    return answer;
}
int bar (int c, int d) {
    int var1 = c + d;
    int var2 = c - d;
    answer = var1 / var2;
    return answer;
}
```

R7
PC after RET

PC during RET

Once bar is finished running and properly torn its stack down to the return value, we are ready to return. Upon calling "RET", the stack looks as it does on the right, the PC is updated with the return value (R7), and the caller, foo, picks it up from there. Bar points R6 at the return value so that foo can load the value immediately using R6.

Address	Value	Description
xEFF6	-4	Callee's Return Value
xEFF7	5	Callee's First param: c
xEFF8	8	Callee's Last param: d
xEFF9	?	Caller's Saved Register(s): Foo saved any registers before use
xEFFA	8	Caller's Local Var(s)
xEFFB	?	Caller's Old Frame Pointer: Foo's caller is main so this is garbage
xEFFC	x3007	Caller's Return Address: Points to main
xEFFD	?	Caller's Return Value: currently unknown
xEFFE	5	Caller's First Param: a
xEFFF	3	Caller's Last Param: b
xF000	?	Empty (main's frame)

Callee's (bar's) stack frame: xEFF6–xEFF8
Caller's (foo's) stack frame: xEFF9–xF000

R6 → xEFF6
R5 → xEFFA

Caller Accepts Return (Transition 4)

All we have to do is load the return value into a register so we can work with it, and then pop the stack pointer back down to the top of our stack.

Tasks

Load return value into a register we want it in
Pop down R6 to top of caller's frame

State 4/0 - During Foo's Execution

```
int foo (int a, int b) {
    int local_var = a + b;
    answer = bar(a, local_var);
    return answer;
}
```

PC

We are now back to where we started! Woo, what a lot of work. Just remember that once you figure out the instructions for transitions 2-4, they can be used as cookie cutters for every function call you ever make. Transitions 1, 2.1, and 3.0 will vary depending on the function (number of params, locals, and regs used). Note: If we wanted to continue, returning out of foo and back to main would look just like returning from bar to foo. Just treat foo as the callee of main.

Address	Value	Description
xEFF9	?	Foo's Saved Register(s): Foo saved all registers it used before use
xEFFA	8	Foo's Local Var(s): stores value of local_var
xEFFB	?	Foo's Old Frame Pointer: Foo's caller is main so this is garbage
xEFFC	x3007	Foo's Return Address: Points to instruction in main to return to
xEFFD	?	Foo's Return Value: currently unknown
xEFFE	5	Foo's First Param: a
xEFFF	3	Foo's Last Param: b
xF000	?	Empty (main's frame)

R6 → xEFF9
R5 → xEFFA

5 Rules and Regulations

5.1 General Rules

1. Starting with the assembly homeworks, any code you write must be meaningfully commented. You should comment your code in terms of the algorithm you are implementing; we all know what each line of code does.

2. Although you may ask TAs for clarification, you are ultimately responsible for what you submit. This means that (in the case of demos) you should come prepared to explain to the TA how any piece of code you submitted works, even if you copied it from the book or read about it on the internet.

3. Please read the assignment in its entirety before asking questions.

4. Please start assignments early, and ask for help early. Do not email us the night the assignment is due with questions.

5. If you find any problems with the assignment it would be greatly appreciated if you reported them to the author (which can be found at the top of the assignment). Announcements will be posted if the assignment changes.

5.2 Submission Conventions

1. All files you submit for assignments in this course should have your name at the top of the file as a comment for any source code file, and somewhere in the file, near the top, for other files unless otherwise noted.

2. When preparing your submission you may either submit the files individually to Canvas/Gradescope or you may submit an archive (zip or tar.gz only please) of the files. You can create an archive by right clicking on files and selecting the appropriate compress option on your system. Both ways (uploading raw files or an archive) are exactly equivalent, so choose whichever is most convenient for you.

3. Do not submit compiled files, that is .class files for Java code and .o files for C code. Only submit the files we ask for in the assignment.

4. Do not submit links to files. The autograder does not understand it, and we will not manually grade assignments submitted this way as it is easy to change the files after the submission period ends.

5.3 Submission Guidelines

1. You are responsible for turning in assignments on time. This includes allowing for unforeseen circumstances. If you have an emergency let us know **IN ADVANCE** of the due time supplying documentation (i.e. note from the dean, doctor's note, etc). Extensions will only be granted to those who contact us in advance of the deadline and no extensions will be made after the due date.

2. You are also responsible for ensuring that what you turned in is what you meant to turn in. After submitting you should be sure to download your submission into a brand new folder and test if it works. No excuses if you submit the wrong files, what you turn in is what we grade. In addition, your assignment must be turned in via Canvas/Gradescope. Under no circumstances whatsoever we will accept any email submission of an assignment. Note: if you were granted an extension you will still turn in the assignment over Canvas/Gradescope.

3. There is a 6-hour grace period added to all assignments. You may submit your assignment without penalty up until 11:55PM, or with 25% penalty up until 5:55AM. So what you should take from this is not to start assignments on the last day and plan to submit right at 11:54AM. You alone are responsible for submitting your homework before the grace period begins or ends; neither Canvas/Gradescope, nor

your flaky internet are to blame if you are unable to submit because you banked on your computer working up until 11:54PM. The penalty for submitting during the grace period (25%) or after (no credit) is non-negotiable.

5.4 Syllabus Excerpt on Academic Misconduct

Academic misconduct is taken very seriously in this class. Quizzes, timed labs and the final examination are individual work.

Homework assignments are collaborative, In addition many if not all homework assignments will be evaluated via demo or code review. During this evaluation, you will be expected to be able to explain every aspect of your submission. Homework assignments will also be examined using computer programs to find evidence of unauthorized collaboration.

What is unauthorized collaboration? Each individual programming assignment should be coded by you. You may work with others, but each student should be turning in their own version of the assignment. Submissions that are essentially identical will receive a zero and will be sent to the Dean of Students' Office of Academic Integrity. Submissions that are copies that have been superficially modified to conceal that they are copies are also considered unauthorized collaboration.

You are expressly forbidden to supply a copy of your homework to another student via electronic means. This includes simply e-mailing it to them so they can look at it. If you supply an electronic copy of your homework to another student and they are charged with copying, you will also be charged. This includes storing your code on any site which would allow other parties to obtain your code such as but not limited to public repositories (Github), pastebin, etc. If you would like to use version control, use github.gatech.edu

5.5 Is collaboration allowed?

Collaboration is allowed on a high level, meaning that you may discuss design points and concepts relevant to the homework with your peers, share algorithms and pseudo-code, as well as help each other debug code. What you shouldn't be doing, however, is pair programming where you collaborate with each other on a single instance of the code. Furthermore, sending an electronic copy of your homework to another student for them to look at and figure out what is wrong with their code is not an acceptable way to help them, because it is frequently the case that the recipient will simply modify the code and submit it as their own. Consider instead using a screen-sharing collaboration app, such as `http://webex.gatech.edu/`, to help someone with debugging if you're not in the same room.

Figure 1: Collaboration rules, explained colorfully